

# Homework 3

Matthew, Ryan, Dani

2017-11-01

## Question 1 (Chapter 6, #8, parts (a)-(e), 10 marks)

a. (1 mark) (Note: You should set your random seed, for reproducibility.)

```
set.seed(1234)
n = 100
X = rnorm(n)
error = rnorm(n)
```

b. (1 mark)

We selected Beta values by the following rule:

$$\beta_i = i + 1$$

```
Y = 1 + 2*X + 3*X^2 +4*X^3 + error
```

c. (3 marks) For the “best model obtained”, you should use one that is parsimonious and close to the consensus best according tht the three selection criteria.

You don't **have** to create a data frame. `regsubsets()` can take a design matrix and response vector, just like `lm.fit()` and `glmnet()`. If you do decide to create a data frame, the following hint may be of use:

```
pmax <- 10
Xmat <- matrix(NA,nrow=n,ncol=pmax)
for(i in 1:pmax) {
  Xmat[,i] <- X^i
}
colnames(Xmat) <- paste0("X.",1:pmax)
dat <- data.frame(Y,Xmat)
```

## Exhaustive Method

```
models = regsubsets(Y ~ ., data = dat, nvmax = 10)
models.sum = summary(models)

cp.best = which.min(models.sum$cp)
bic.best = which.min(models.sum$bic)
rsq.best = which.max(models.sum$adjr2)

print(paste("Best model for Cp is model with", cp.best, "variables"))
```

```
## [1] "Best model for Cp is model with 3 variables"
```

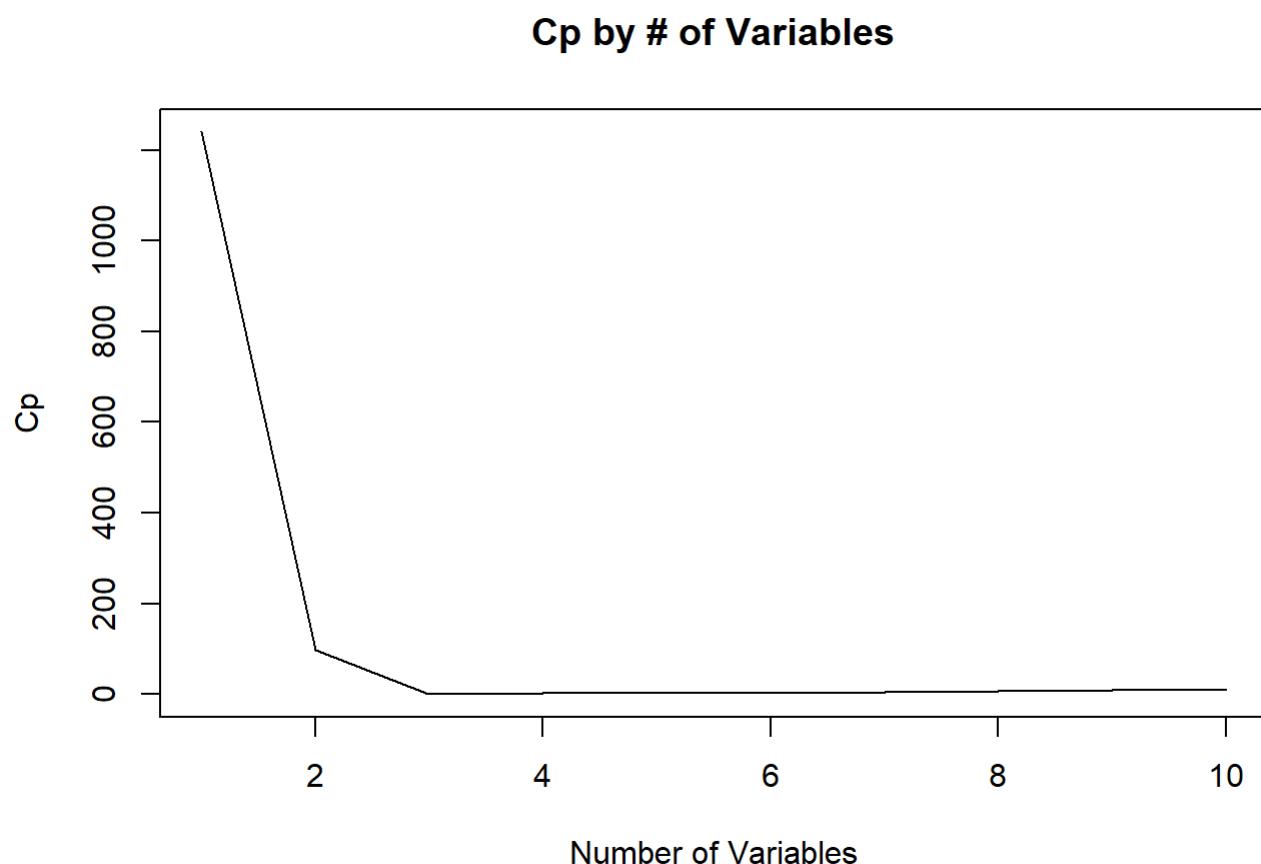
```
print(paste("Best model for BIC is model with", bic.best, "variables"))
```

```
## [1] "Best model for BIC is model with 3 variables"
```

```
print(paste("Best model for Adjusted R^2 is model with", rsq.best, "variables"))
```

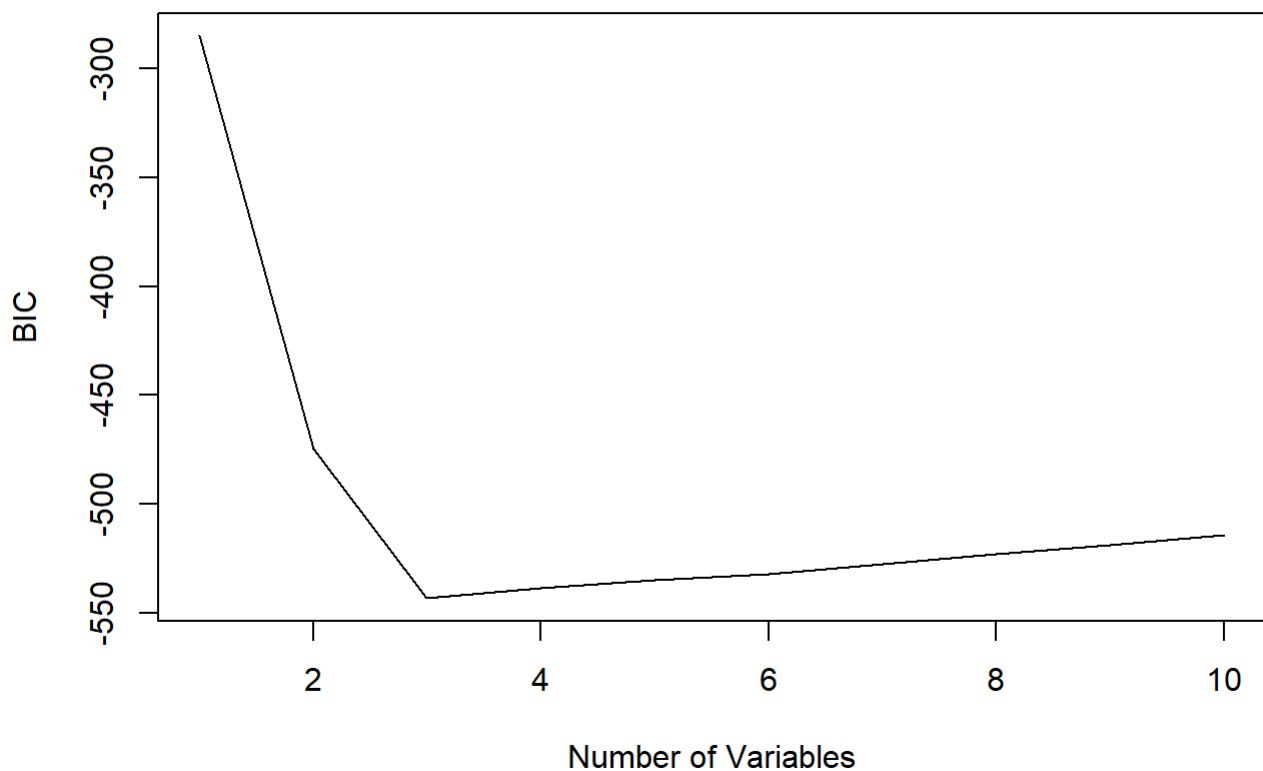
```
## [1] "Best model for Adjusted R^2 is model with 3 variables"
```

```
plot(models.sum$cp, xlab = "Number of Variables", ylab = "Cp", main = "Cp by # of Variables", type="l")
```



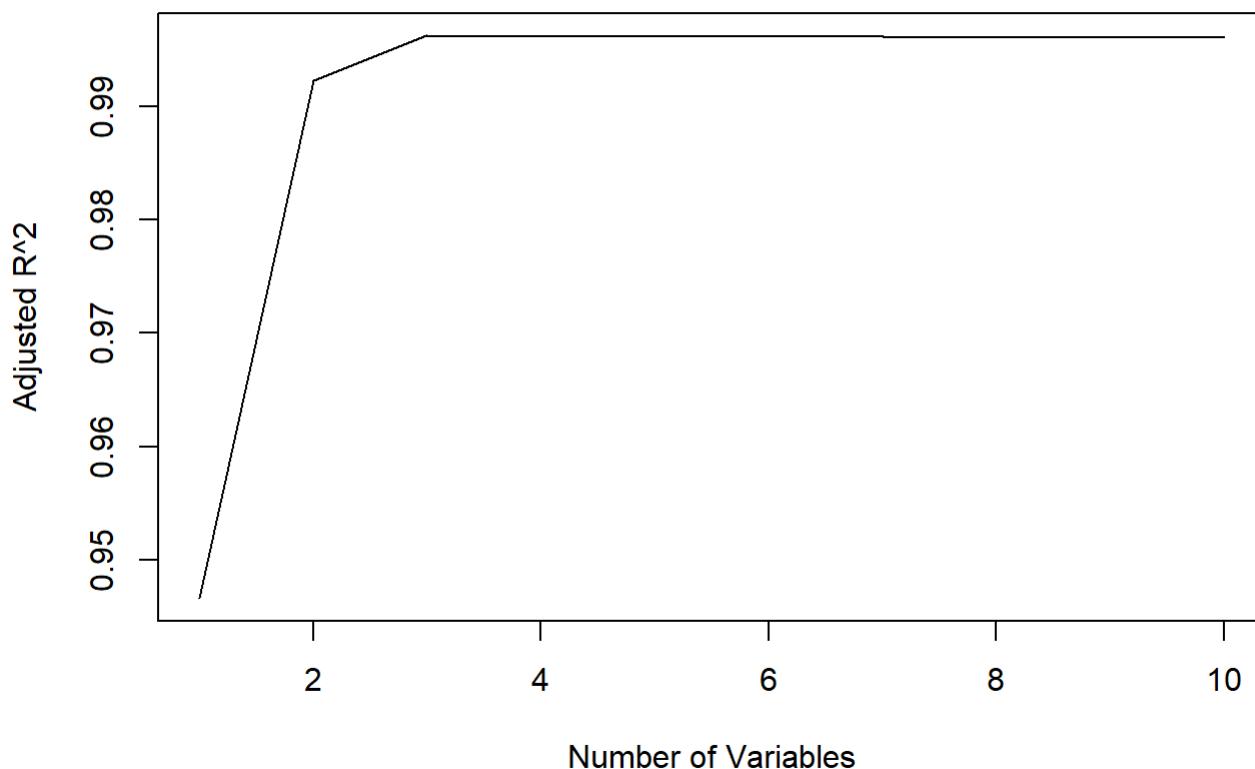
```
plot(models.sum$bic, xlab = "Number of Variables", ylab = "BIC", main = "BIC by # of Variables", type="l")
```

## BIC by # of Variables



```
plot(models.sum$adjr2, xlab = "Number of Variables", ylab = "Adjusted R^2", main = "Adjusted R^2  
by # of Variables", type="l")
```

### Adjusted R<sup>2</sup> by # of Variables



```
if(cp.best == bic.best & bic.best == rsq.best){  
  print("Best model for all variables has coefficients")  
  coef(models, cp.best)  
} else{  
  print("Best model for Cp has coefficients")  
  coef(models, cp.best)  
  
  print("Best model for BIC has coefficients")  
  coef(models, bic.best)  
  
  print("Best model for Adjusted R^2 has coefficients")  
  coef(models, rsq.best)  
}
```

```
## [1] "Best model for all variables has coefficients"
```

```
## (Intercept)      X.1      X.2      X.3  
## 1.132470    1.912586   2.893627   4.032305
```

d. (2 marks)

## Forward Selection

```

models = regsubsets(Y ~ ., method = "forward", data = dat, nvmax = 10)
models.sum = summary(models)

cp.best = which.min(models.sum$cp)
bic.best = which.min(models.sum$bic)
rsq.best = which.max(models.sum$adjr2)

print(paste("Best model for Cp is model with", cp.best, "variables"))

## [1] "Best model for Cp is model with 3 variables"

print(paste("Best model for BIC is model with", bic.best, "variables"))

## [1] "Best model for BIC is model with 3 variables"

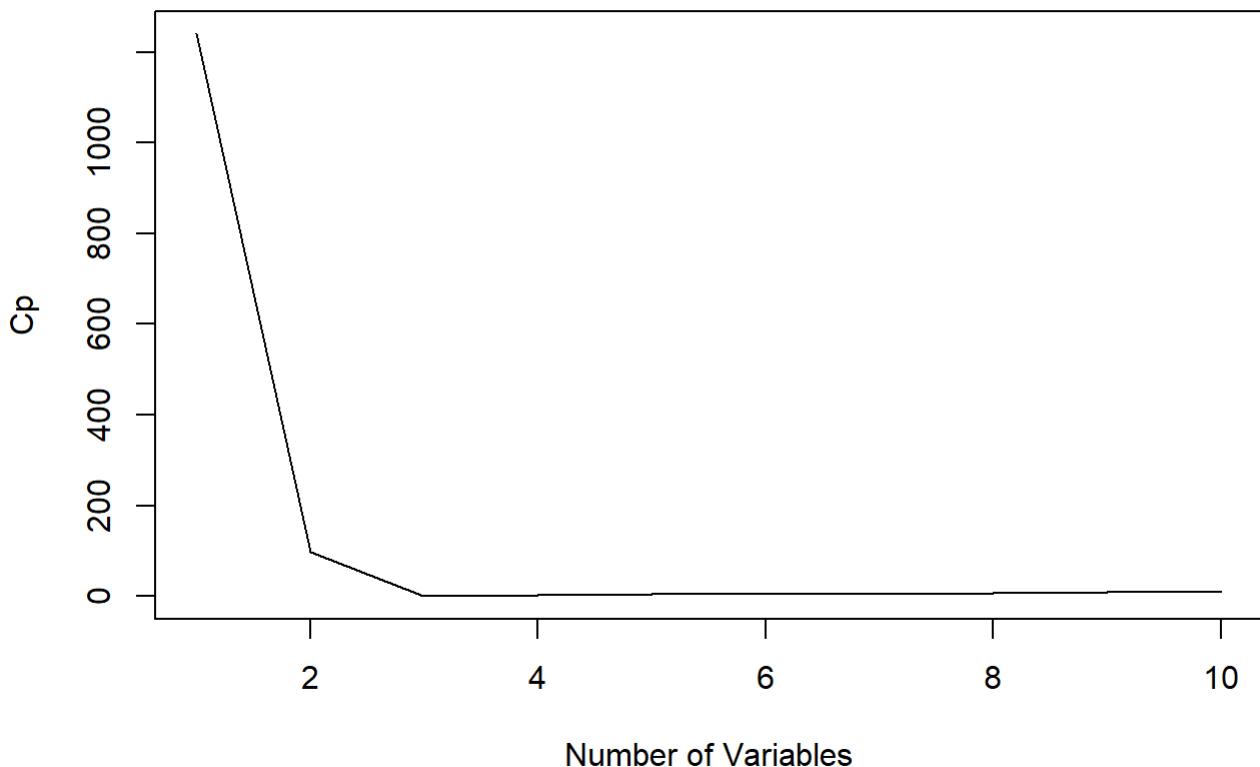
print(paste("Best model for Adjusted R^2 is model with", rsq.best, "variables"))

## [1] "Best model for Adjusted R^2 is model with 3 variables"

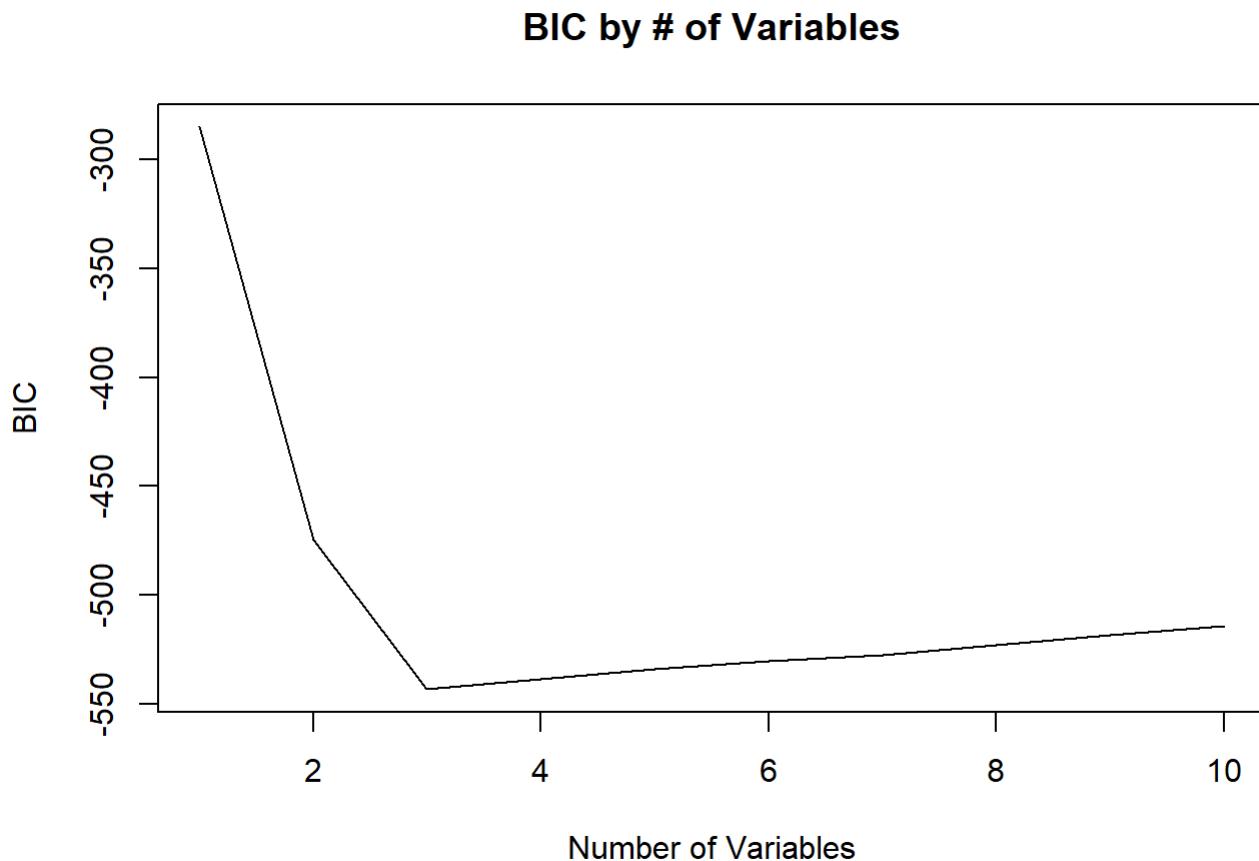
plot(models.sum$cp, xlab = "Number of Variables", ylab = "Cp", main = "Cp by # of Variables", type="l")

```

**Cp by # of Variables**

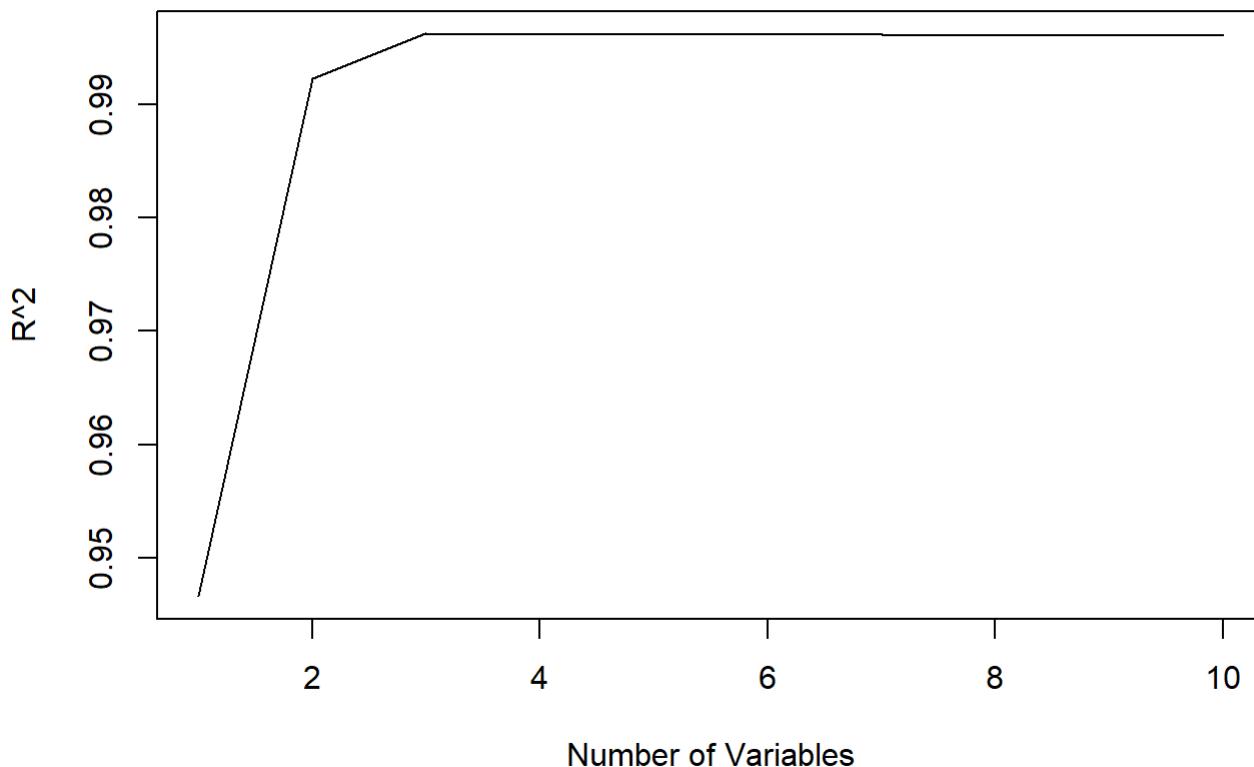


```
plot(models.sum$bic, xlab = "Number of Variables", ylab = "BIC", main = "BIC by # of Variables", type="l")
```



```
plot(models.sum$adjr2, xlab = "Number of Variables", ylab = "R^2", main = "Adjusted R^2 by # of Variables", type="l")
```

## Adjusted R<sup>2</sup> by # of Variables



```
if(cp.best == bic.best & bic.best == rsq.best){  
  print("Best model for all variables has coefficients")  
  coef(models, cp.best)  
} else{  
  print("Best model for Cp has coefficients")  
  coef(models, cp.best)  
  
  print("Best model for BIC has coefficients")  
  coef(models, bic.best)  
  
  print("Best model for Adjusted R^2 has coefficients")  
  coef(models, rsq.best)  
}
```

```
## [1] "Best model for all variables has coefficients"
```

```
## (Intercept)      X.1      X.2      X.3  
## 1.132470    1.912586   2.893627   4.032305
```

## Backward Selection

```

models = regsubsets(Y ~ ., method = "backward", data = dat, nvmax = 10)
models.sum = summary(models)

cp.best = which.min(models.sum$cp)
bic.best = which.min(models.sum$bic)
rsq.best = which.max(models.sum$adjr2)

print(paste("Best model for Cp is model with", cp.best, "variables"))

## [1] "Best model for Cp is model with 3 variables"

print(paste("Best model for BIC is model with", bic.best, "variables"))

## [1] "Best model for BIC is model with 3 variables"

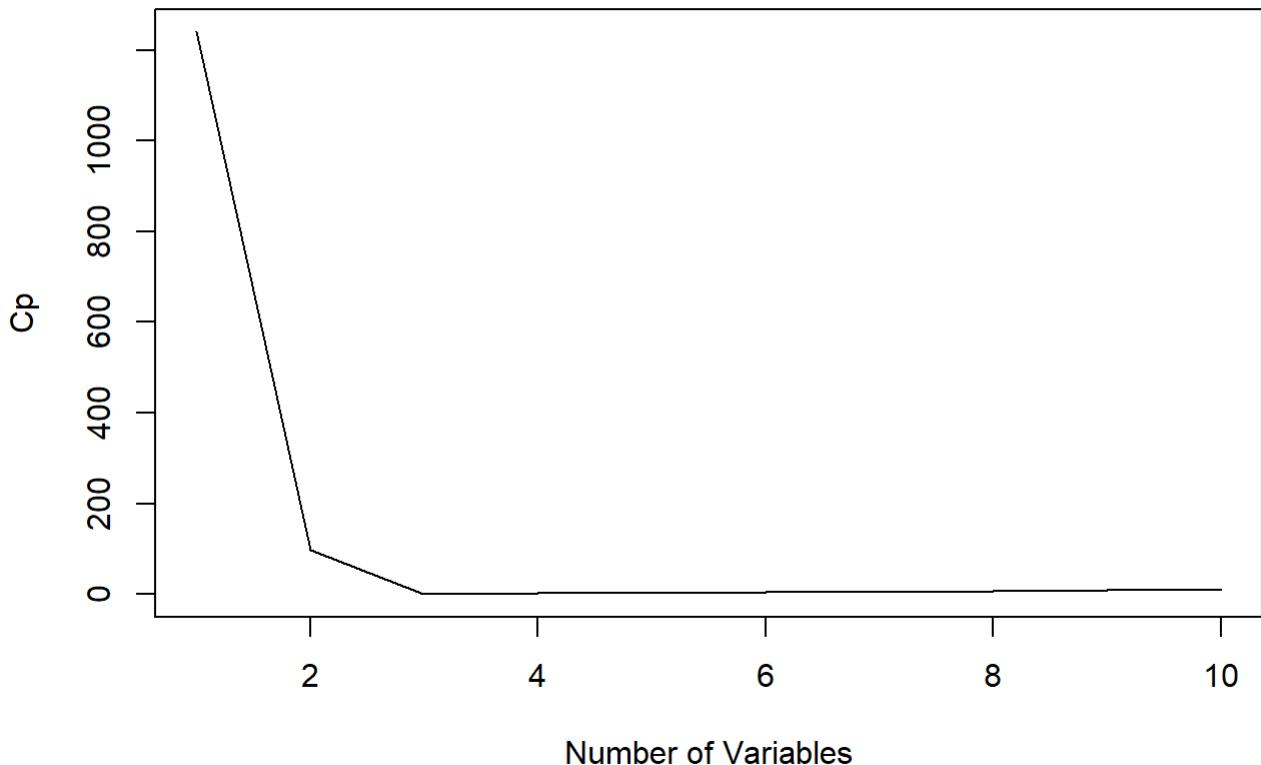
print(paste("Best model for Adjusted R^2 is model with", rsq.best, "variables"))

## [1] "Best model for Adjusted R^2 is model with 3 variables"

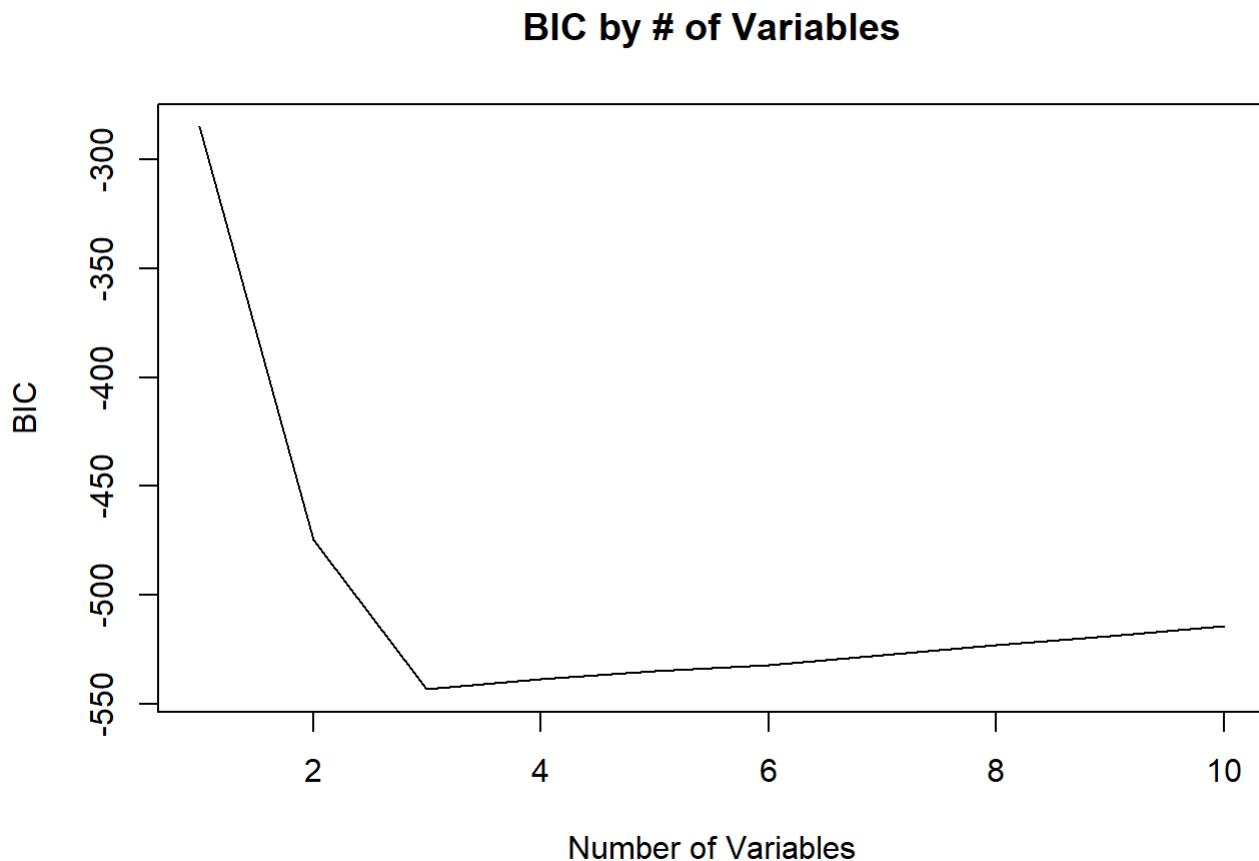
plot(models.sum$cp, xlab = "Number of Variables", ylab = "Cp", main = "Cp by # of Variables", type="l")

```

**Cp by # of Variables**

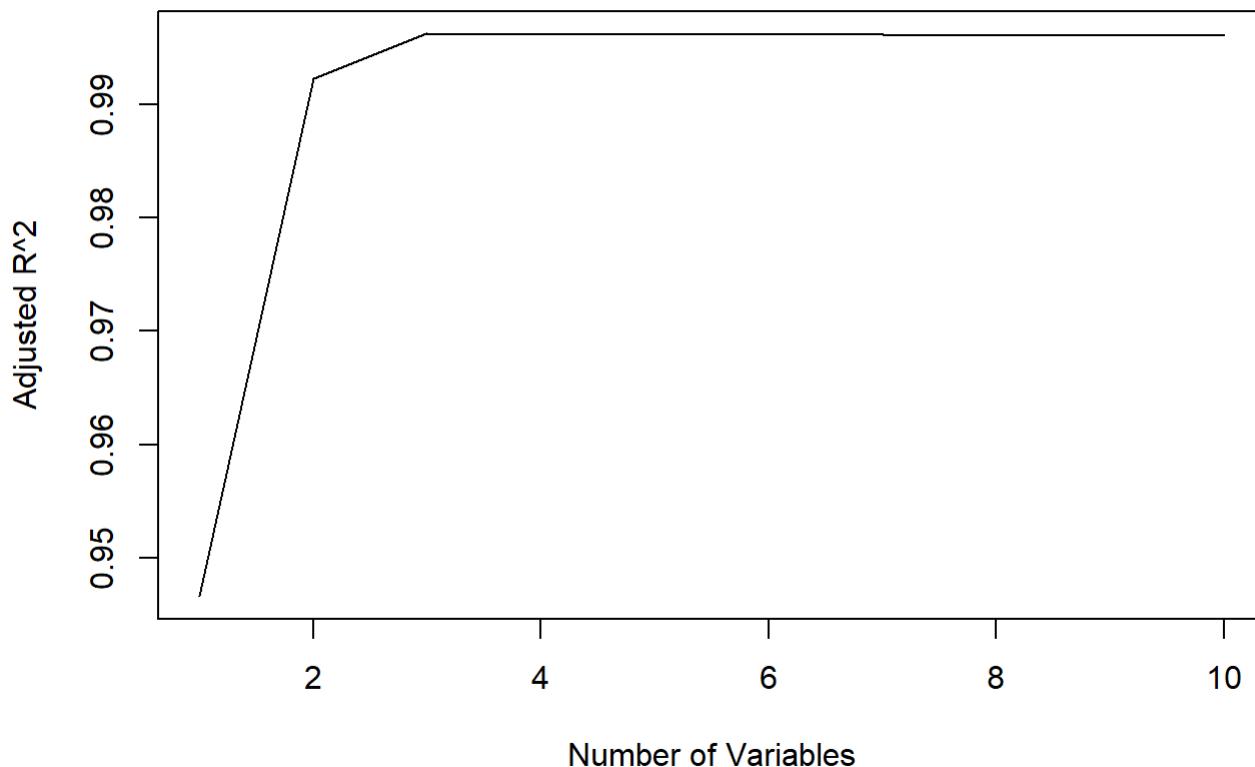


```
plot(models.sum$bic, xlab = "Number of Variables", ylab = "BIC", main = "BIC by # of Variables", type="l")
```



```
plot(models.sum$adjr2, xlab = "Number of Variables", ylab = "Adjusted R^2", main = "Adjusted R^2 by # of Variables", type="l")
```

## Adjusted R<sup>2</sup> by # of Variables



```
if(cp.best == bic.best & bic.best == rsq.best){  
  print("Best model for all variables has coefficients")  
  coef(models, cp.best)  
} else{  
  print("Best model for Cp has coefficients")  
  coef(models, cp.best)  
  
  print("Best model for BIC has coefficients")  
  coef(models, bic.best)  
  
  print("Best model for Adjusted R^2 has coefficients")  
  coef(models, rsq.best)  
}
```

```
## [1] "Best model for all variables has coefficients"
```

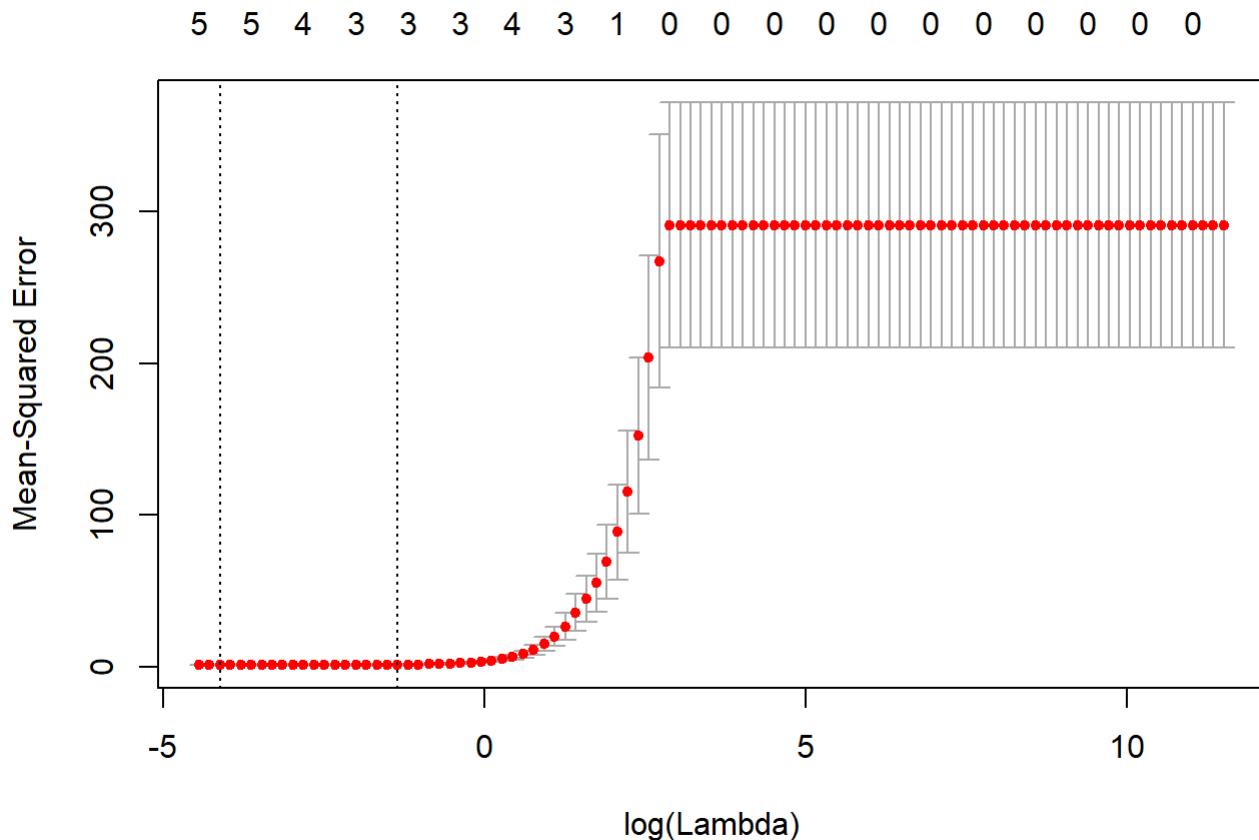
```
## (Intercept)      X.1      X.2      X.3  
## 1.132470    1.912586   2.893627   4.032305
```

## Compare to results in 8c

The results are exactly the same for all 3 selection criteria. The same model was selected by all criteria in all directions.

e. (3 marks)

```
lambdas <- 10^{seq(from=-2,to=5,length=100)}
cv.lafit <- cv.glmnet(Xmat,Y,alpha=1,lambda=lambdas)
plot(cv.lafit)
```



```
la.best.lam <- cv.lafit$lambda.1se
la.best.lam
```

```
## [1] 0.2595024
```

```
la.best <- glmnet(Xmat,Y,alpha=1,lambda=la.best.lam)
coef(la.best)
```

```

## 11 x 1 sparse Matrix of class "dgCMatrix"
##           s0
## (Intercept) 1.268441
## X.1         1.754901
## X.2         2.738949
## X.3         4.012495
## X.4         .
## X.5         .
## X.6         .
## X.7         .
## X.8         .
## X.9         .
## X.10        .

```

## Discuss results obtained from CV and lasso model

The results are almost identical to the model selected above. The coefficients of the predictors in the lasso model have seen a very slight shrinkage the intercept has seen a small increase.

### Question 2 (Ch6, #9, 12 marks)

a. (0 marks) To make everyone's results comparable, please select your test set with the following.

```

data(College)
# Standardize columns
College <- mutate(College, Private = as.numeric(Private=="Yes"))
College <- data.frame(lapply(College,scale))
dim(College) # 777 rows, use 111 as test

```

```
## [1] 777 18
```

```

set.seed(1)
testset <- sample(1:777,size=111)
College.test <- College[testset,]
College.train <- College[-testset,]

```

b. (2 marks)

```

linear = lm(Apps ~ . , data = College.train)

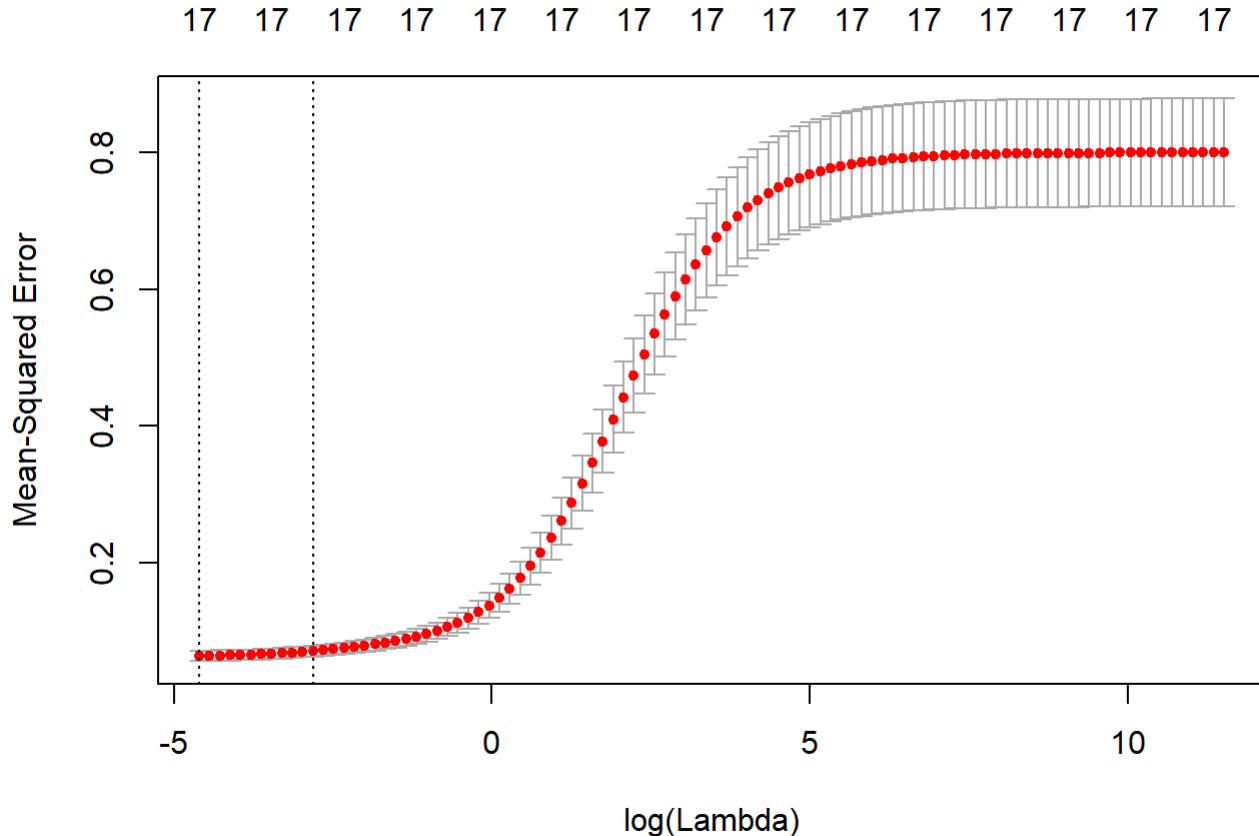
preds = predict(linear, College.test)
e1 = mean((College.test$Apps - preds)^2)
e1

```

```
## [1] 0.1866961
```

c. (2 marks)

```
lambdas <- 10^{seq(from=-2,to=5,length=100)}
X = College.train %>% select(-Apps) %>% data.matrix()
y = College.train$Apps
cv.lafit <- cv.glmnet(X, y, alpha=0, lambda=lambdas)
plot(cv.lafit)
```



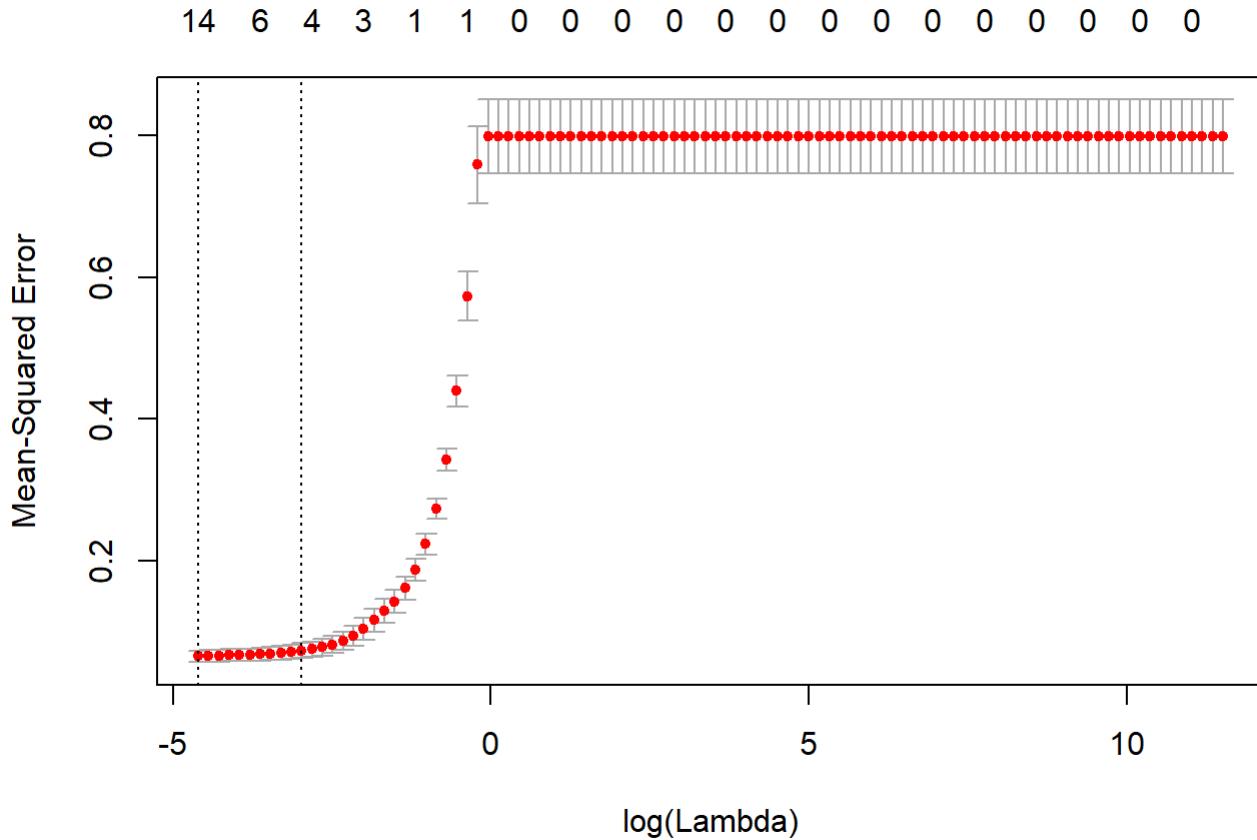
```
la.best.lam <- cv.lafit$lambda.1se
la.best <- glmnet(X, y, alpha=0, lambda = la.best.lam)

x.pred = College.test %>% select(-Apps) %>% data.matrix()
preds = predict(la.best, x.pred)
e2 = mean((College.test$Apps - preds)^2)
e2
```

```
## [1] 0.3363783
```

d. (2 marks)

```
lambdas <- 10^{seq(from=-2,to=5,length=100)}
X = College.train %>% select(-Apps) %>% data.matrix()
y = College.train$Apps
cv.lafit <- cv.glmnet(X, y, alpha = 1, lambda = lambdas)
plot(cv.lafit)
```



```

la.best.lam <- cv.lafit$lambda.1se
la.best <- glmnet(X, y, alpha = 1, lambda = la.best.lam)

x.pred = College.test %>% select(-Apps) %>% data.matrix()
preds = predict(la.best, x.pred)
e3 = mean((College.test$Apps - preds)^2)

print(paste("number of coefs larger than 0:", sum(coef(la.best) > 0)))

```

```
## [1] "number of coefs larger than 0: 4"
```

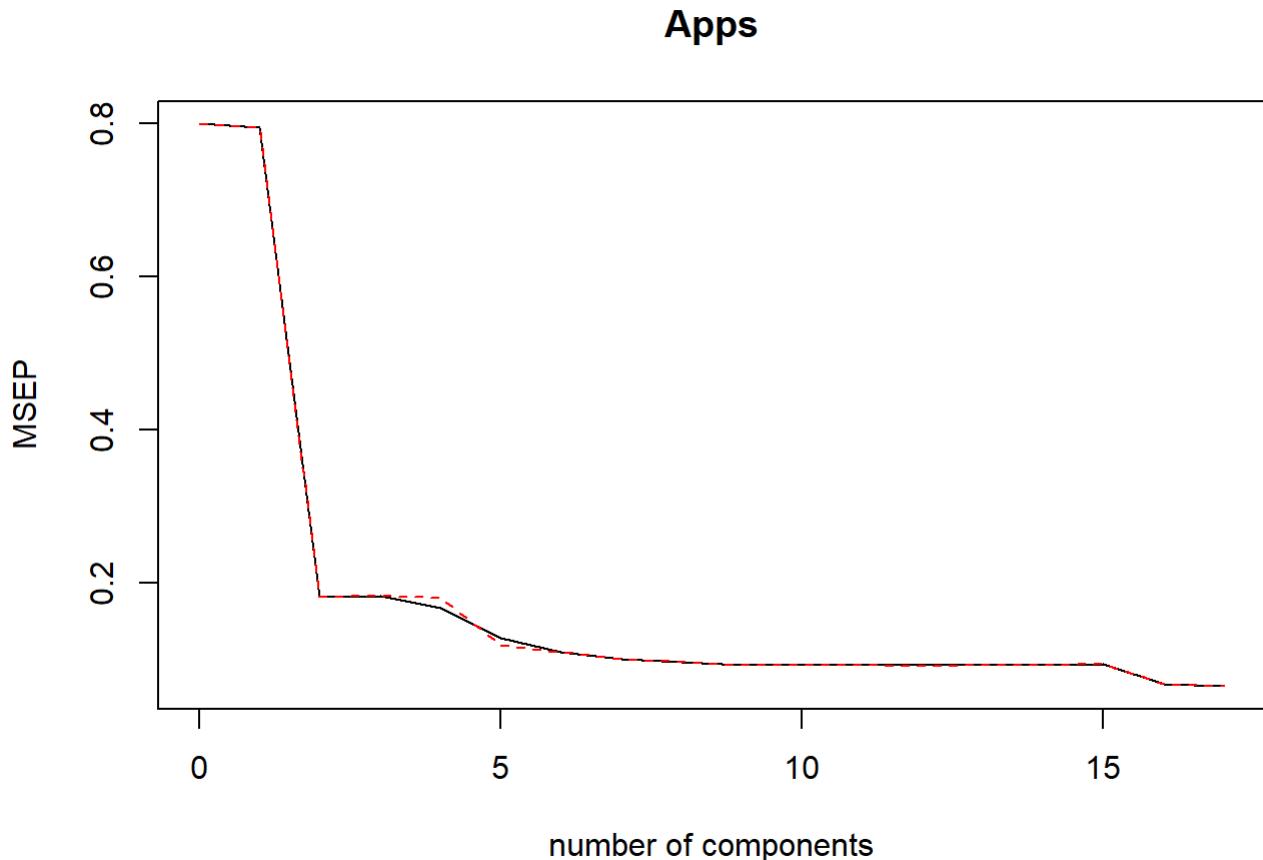
e3

```
## [1] 0.2490143
```

e. (2 marks)

```
# data already scaled above
pcr.fit = pcr(Apps ~ ., data = College.train, validation = "CV")

# Choose M by the graph
validationplot(pcr.fit, val.type = "MSEP")
```



```
print("We choose a value of 5 for M is a simple model that has a very good % variance explained")
```

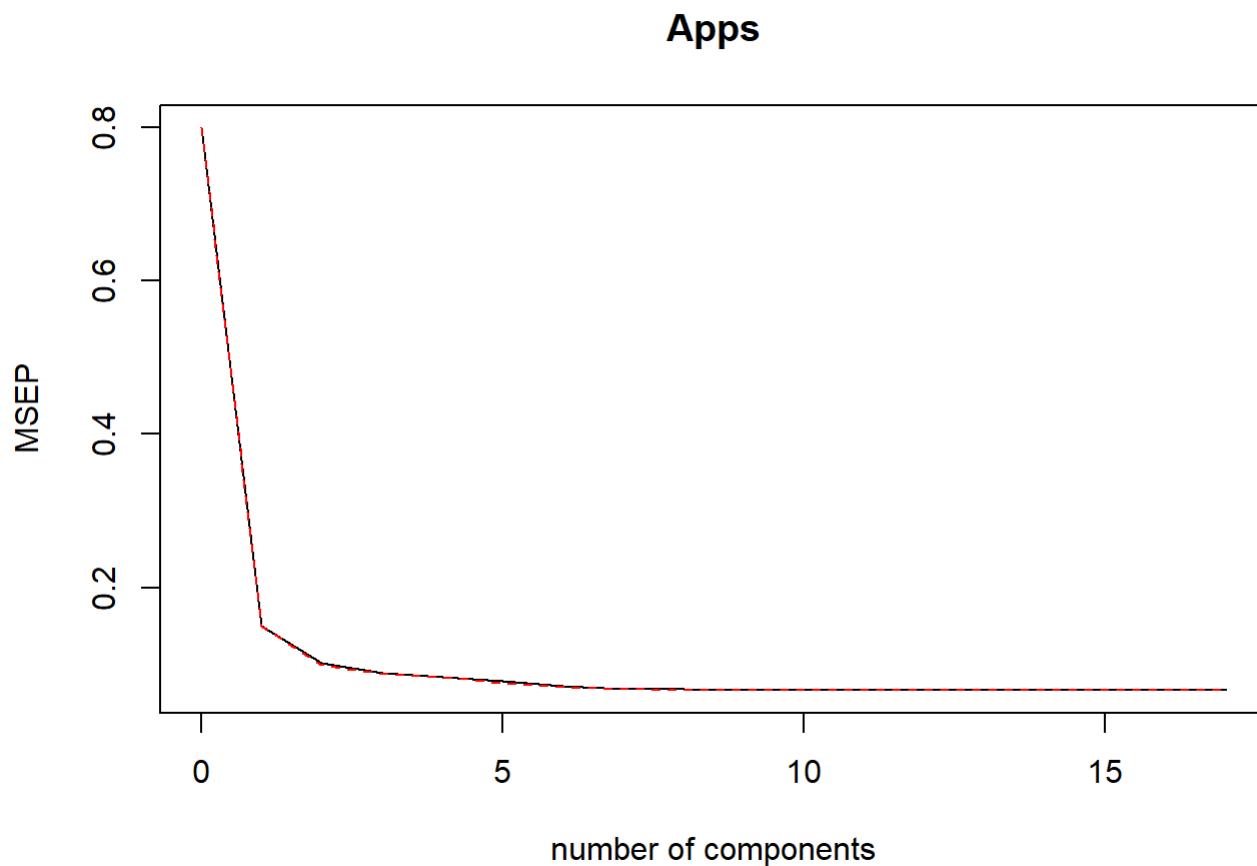
```
## [1] "We choose a value of 5 for M is a simple model that has a very good % variance explained"
```

```
preds = predict(pcr.fit, College.test, ncomp = 5)
e4 = mean((preds - College.test$Apps)^2)
e4
```

```
## [1] 0.5938384
```

f. (2 marks)

```
pls.fit = plsr(Apps ~ ., data = College.train, validation = "CV")
# Choose M by the graph and summary
validationplot(pls.fit, val.type = "MSEP")
```



```
summary(pls.fit)
```

```

## Data: X dimension: 666 17
## Y dimension: 666 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV          0.8943  0.3874  0.3181  0.2964  0.2897  0.2798  0.2669
## adjCV       0.8943  0.3867  0.3147  0.2959  0.2882  0.2752  0.2653
##      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV          0.2611  0.2591  0.2585  0.2582  0.2579  0.2583  0.2579
## adjCV       0.2601  0.2583  0.2577  0.2573  0.2570  0.2574  0.2570
##      14 comps 15 comps 16 comps 17 comps
## CV          0.2580  0.2580   0.258   0.2580
## adjCV       0.2571  0.2571   0.257   0.2571
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X          25.12   36.86   62.35   65.59   67.94   73.39   76.29
## Apps       81.70   88.15   89.58   90.66   92.05   92.42   92.57
##      8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          80.57   82.82   84.84   88.40   91.49   93.03   95.54
## Apps       92.60   92.63   92.67   92.69   92.69   92.70   92.70
##      15 comps 16 comps 17 comps
## X          97.18   98.96   100.0
## Apps       92.70   92.70   92.7

```

```

print("We choose a value of 2 for M as it is the simplest model that has a acceptable error rate")

```

```

## [1] "We choose a value of 2 for M as it is the simplest model that has a acceptable error rate"

```

```

preds = predict(pls.fit, College.test, ncomp = 2)
e5 = mean((preds - College.test$Apps)^2)
e5

```

```

## [1] 0.5186612

```

g. (2 marks)

How accurately can we predict the number of college applications received?

Is there much difference among the test errors resulting from these five approaches?

```

e1 #Lm

```

```

## [1] 0.1866961

```

```
e3 # Lasso
```

```
## [1] 0.2490143
```

```
e2 # ridge
```

```
## [1] 0.3363783
```

```
e5 # pls
```

```
## [1] 0.5186612
```

```
e4 # pcr
```

```
## [1] 0.5938384
```

We can predict the number of college applications very closely. The MSE is very small compared to the magnitude of the response variable. Our in sample Adjusted R-squared is 0.92.

The linear model has the smallest MSE for the testing data. It is significantly smaller than that of pcr. There is a spread in MSE across the 5 methods. This indicates that there is a distinct ranking between the 5 methods.

## Question 3 (Ch7, #6, 8 marks)

a. (5 marks)

```

attach(Wage)
set.seed(444)

# Make sure train set has full range of age values
Wage <- Wage %>% arrange(age)
Wage.train = Wage[c(1,nrow(Wage)), ]
Wage2 <- Wage[-c(1, nrow(Wage)), ]

# make test set and train set
testset <- sample(1:2998, size = 300)
Wage.test <- Wage2[testset,]
Wage.train <- rbind(Wage2[-testset,], Wage.train)

errors = rep(NA, 8)
for(i in 1:8){
  fit = lm(wage ~ poly(age, i), data = Wage.test)
  preds = predict(fit, Wage.train)
  errors[i] = mean((Wage.train$wage - preds)^2)
}

d = which.min(errors)

fit.1=lm(wage~age,data=Wage)
fit.2=lm(wage~poly(age, 2), data=Wage)
fit.3=lm(wage~poly(age, 3), data=Wage)
fit.4=lm(wage~poly(age, 4), data=Wage)
fit.5=lm(wage~poly(age, 5), data=Wage)
fit.6=lm(wage~poly(age, 6), data=Wage)
fit.7=lm(wage~poly(age, 7), data=Wage)
fit.8=lm(wage~poly(age, 8), data=Wage)
anova(fit.1, fit.2, fit.3, fit.4, fit.5, fit.6, fit.7, fit.8)

```

```

## Analysis of Variance Table
##
## Model 1: wage ~ age
## Model 2: wage ~ poly(age, 2)
## Model 3: wage ~ poly(age, 3)
## Model 4: wage ~ poly(age, 4)
## Model 5: wage ~ poly(age, 5)
## Model 6: wage ~ poly(age, 6)
## Model 7: wage ~ poly(age, 7)
## Model 8: wage ~ poly(age, 8)
##   Res.Df   RSS Df Sum of Sq      F    Pr(>F)
## 1 2998 5022216
## 2 2997 4793430  1   228786 143.6484 < 2.2e-16 ***
## 3 2996 4777674  1    15756  9.8926  0.001676 **
## 4 2995 4771604  1     6070  3.8113  0.051002 .
## 5 2994 4770322  1     1283  0.8053  0.369590
## 6 2993 4766389  1     3932  2.4690  0.116221
## 7 2992 4763834  1     2555  1.6044  0.205381
## 8 2991 4763707  1     127  0.0795  0.777952
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

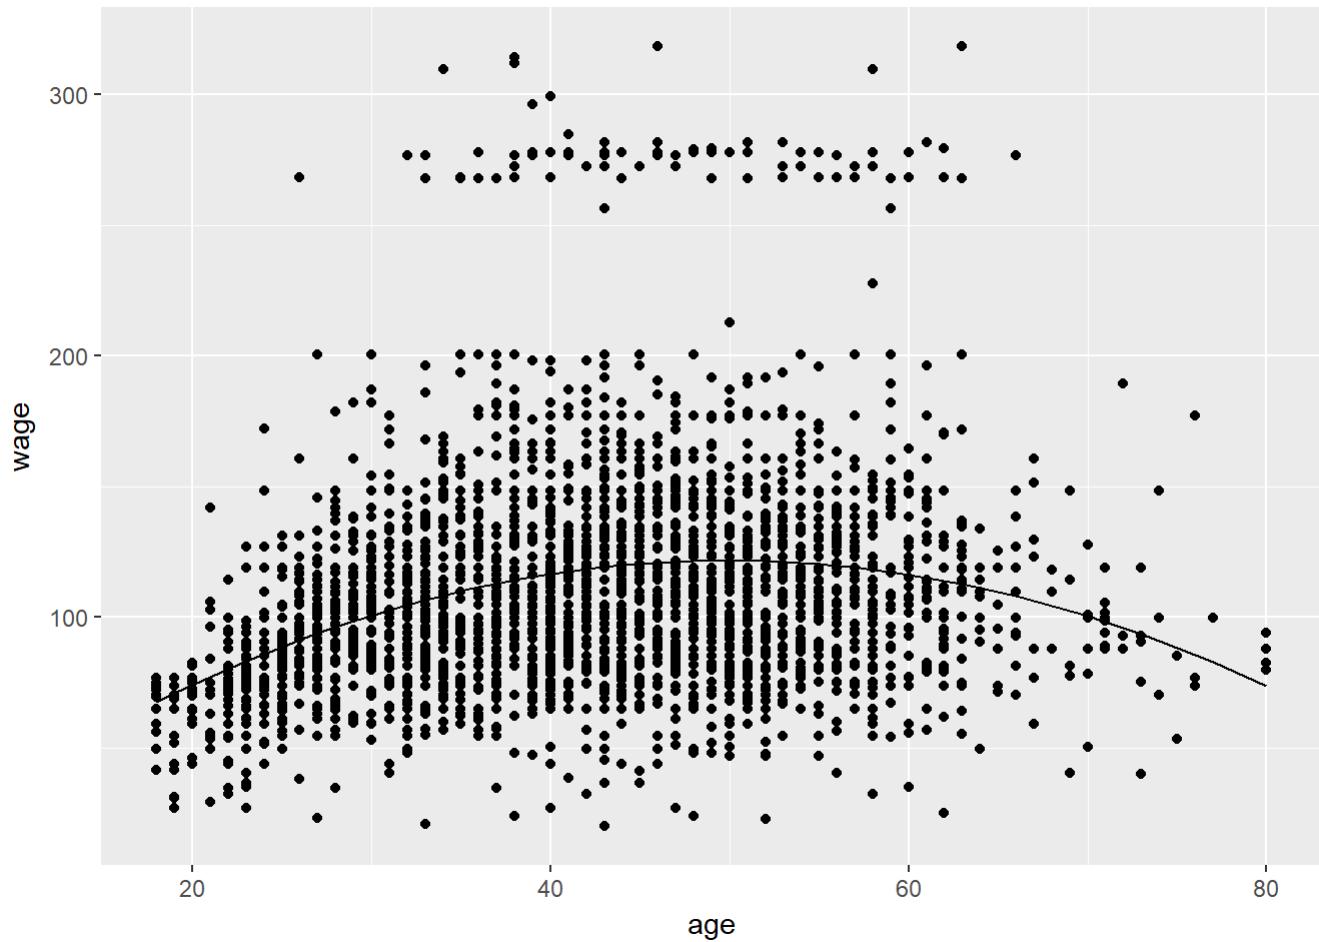
```

```

dat = data.frame(Wage$age, fitted(fit.2))

ggplot() +
  geom_point(data = Wage, aes(x = age, y = wage)) +
  geom_line(data = dat, aes(x = Wage.age, y = fitted.fit.2.))

```



The degree chose was a polynomial of degree 2. This is confirmed by our anova analysis.

b. (3 marks)

```

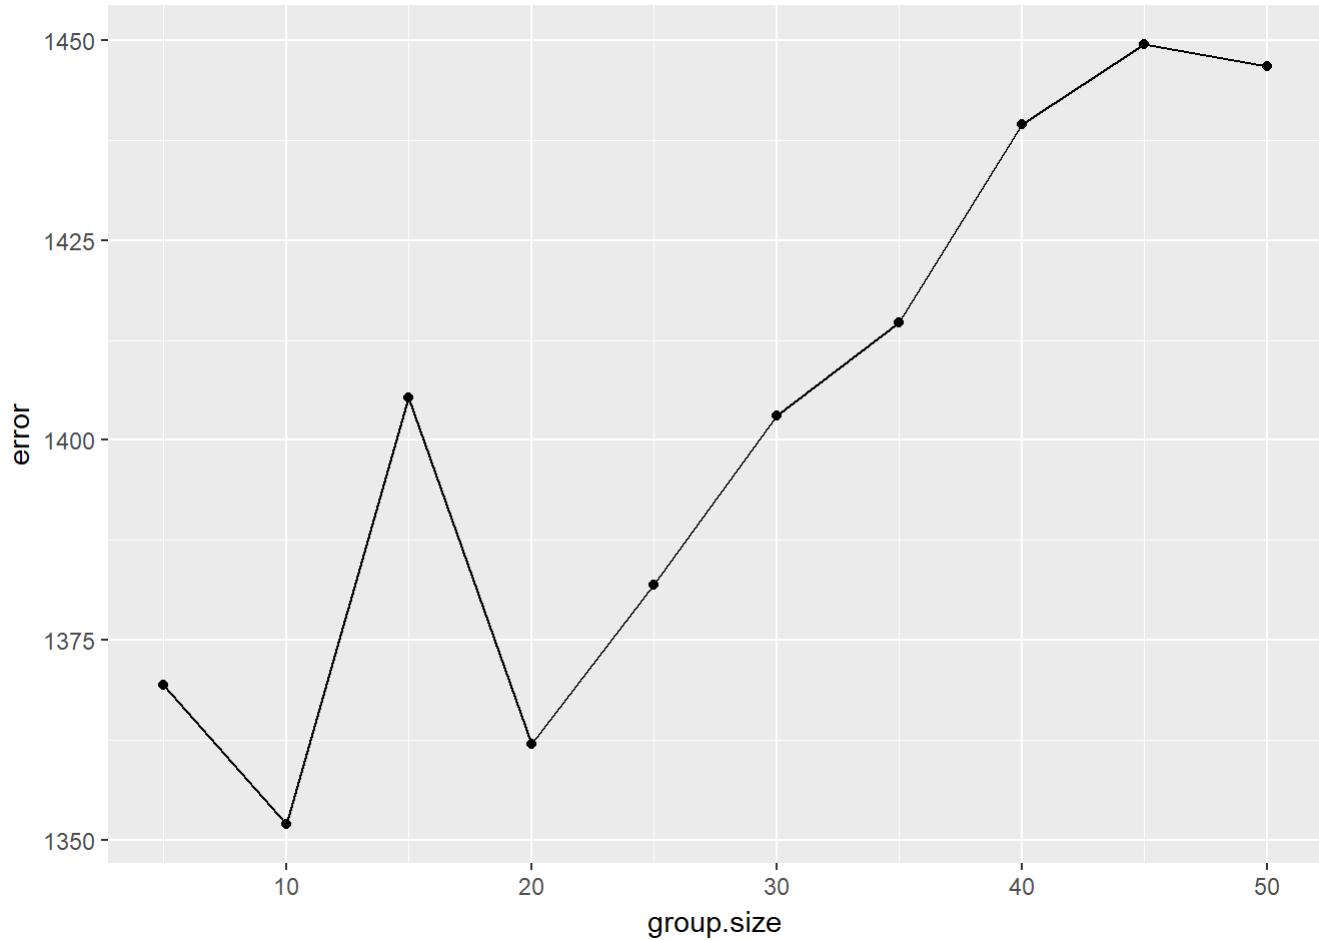
errors = rep(NA, 10)
for(i in 1:10){
  j = i*5
  agebreaks <- seq(17, 80+j, by = j)
  fit = lm(wage ~ cut(age, breaks = agebreaks), data = Wage.train)
  preds <- predict(fit, Wage.test)
  errors[i] <- mean((Wage.test$wage - preds)^2)
}

df.errors = data.frame(seq(5,50, by = 5), errors)
colnames(df.errors) = c("group.size","error")

ggplot(df.errors, aes(x = group.size, y = error))+  

  geom_point()+
  geom_line()

```



```
j = which.min(errors)*5

# age groups should be of size 10
j
```

```
## [1] 10
```

```
agebreaks <- seq(17, 80+j, by = j)
fit = lm(wage ~ cut(age, breaks = agebreaks), data = Wage)
uniqueAge <- data.frame(age = sort(unique(Wage$age)))
preds <- data.frame(uniqueAge,predict(fit,newdata = uniqueAge,interval="confidence"))

ggplot(Wage,aes(x=age,y=wage)) +
  geom_point(alpha=0.1) +
  geom_ribbon(aes(x=age,y=fit,ymin=lwr,ymax=upr), data=preds, fill="blue", alpha=.2) +
  geom_line(aes(y = fit), data = preds, color="blue")
```

